

title: "Homework 3" author: "Arjun Laxman" toc: true title-block-banner: true title-block-style: default
format: pdf # format: pdf

Important

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to only upload a PDF file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

⚠ Please add your name to the author information in the frontmatter before submitting your assignment ⚠

For this assignment, we will be using the [Wine Quality](#) dataset from the UCI Machine Learning Repository. The dataset consists of red and white *vinho verde* wine samples, from the north of Portugal. The goal is to model wine quality based on physicochemical tests

We will be using the following libraries:

```
options(repos = c(CRAN = "https://cloud.r-project.org"))  
  
install.packages("car")
```

The downloaded binary packages are in

/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("corrplot")
```

The downloaded binary packages are in

/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("tidyverse")
```

The downloaded binary packages are in

/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("dplyr")
```

The downloaded binary packages are in

/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("tidyverse")
```

The downloaded binary packages are in
/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("glmnet")
```

The downloaded binary packages are in
/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
install.packages("curl")
```

The downloaded binary packages are in
/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
library(MASS)  
install.packages("glmnet")
```

The downloaded binary packages are in
/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
library(readr)  
library(tidyr)  
library(corrplot)
```

corrplot 0.92 loaded

```
library(dplyr)
```

Attaching package: 'dplyr'

The following object is masked from 'package:MASS':

select

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(purrr)
library(car)
```

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:purrr':

some

The following object is masked from 'package:dplyr':

recode

```
library(glmnet)
```

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8

```
install.packages("curl")
```

The downloaded binary packages are in

/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpW3LrSI/downloaded_packages

```
make_formula <- function(x){
  as.formula(
    paste("quality ~ ", paste(x, collapse = " + "))
  )
}

# Store the formula in a variable
my_formula <- make_formula(c("a", "b", "c"))

# Print the formula
print(my_formula)
```

quality ~ a + b + c

<environment: 0x11db30710>

```
# generate a model matrix for glmnet()
make_model_matrix <- function(formula){
  X <- model.matrix(formula, df)[, -1]
  cnames <- colnames(X)
  for(i in 1:ncol(X)){
    if(!cnames[i] == "typewhite"){
      X[, i] <- scale(X[, i])
    } else {
      colnames(X)[i] <- "type"
    }
  }
  return(X)
}
```

Question 1

50 points

Regression with categorical covariate and *t*-Test

1.1 (5 points)

Read the wine quality datasets from the specified URLs and store them in data frames `df1` and `df2`.

```
url1 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv"
url2 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"

df1 <- read_delim(url1, delim=";")
```

Rows: 4898 Columns: 12

— Column specification —

Delimiter: ";"

dbl (12): fixed acidity, volatile acidity, citric acid, residual sugar, chlo...

i Use ``spec()`` to retrieve the full column specification for this data.

i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
df2 <- read_delim(url2, delim=";")
```

Rows: 1599 Columns: 12

— Column specification —

Delimiter: ";"

dbl (12): fixed acidity, volatile acidity, citric acid, residual sugar, chlo...

- i Use ``spec()`` to retrieve the full column specification for this data.
- i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

1.2 (5 points)

Perform the following tasks to prepare the data frame `df` for analysis:

1. Combine the two data frames into a single data frame `df`, adding a new column called `type` to indicate whether each row corresponds to white or red wine.
2. Rename the columns of `df` to replace spaces with underscores
3. Remove the columns `fixed_acidity` and `free_sulfur_dioxide`
4. Convert the `type` column to a factor
5. Remove rows (if any) with missing values.

```
library(dplyr)
library(stringr) # Ensure this library is loaded

df <- bind_rows(
  df1 %>% mutate(type = "white"),
  df2 %>% mutate(type = "red")
) %>%
  rename_all(~str_replace_all(., " ", "_")) %>%
  select(-c(fixed_acidity, free_sulfur_dioxide)) %>%
  mutate(type = as.factor(type)) %>%
  drop_na()
```

Your output to `R dim(df)` should be

```
[1] 6497  11
```

1.3 (20 points)

Recall from STAT 200, the method to compute the t statistic for the the difference in means (with the equal variance assumption)

1. Using `df` compute the mean of `quality` for red and white wine separately, and then store the difference in means as a variable called `diff_mean`.
2. Compute the pooled sample variance and store the value as a variable called `sp_squared`.
3. Using `sp_squared` and `diff_mean`, compute the t Statistic, and store its value in a variable called `t1`.

```
df_summary <- df %>%
  group_by(type) %>%
  summarize(
    mean = mean(quality),
    sd = sd(quality),
    n = length(quality)
  )
```

```
diff_mean <- df_summary$mean %>% diff()
sp <- sum(
  df_summary$sd^2 * (df_summary$n - 1)
) / sum(df_summary$n - 2)

t1 <- diff_mean / (sqrt(sp) * sqrt(1/nrow(df1) + 1/nrow(df2)))
```

1.4 (10 points)

Equivalently, R has a function called `t.test()` which enables you to perform a two-sample t -Test without having to compute the pooled variance and difference in means.

Perform a two-sample t -test to compare the quality of white and red wines using the `t.test()` function with the setting `var.equal=TRUE`. Store the t -statistic in `t2`.

```
t_test <- t.test(
  df %>% filter(type == "white") %>% select(quality),
  df %>% filter(type == "red") %>% select(quality),
  var.equal=TRUE
)
t2 <- t_test$statistic
```

1.5 (5 points)

Fit a linear regression model to predict `quality` from `type` using the `lm()` function, and extract the t -statistic for the `type` coefficient from the model summary. Store this t -statistic in `t3`.

```
fit <- lm(quality ~ type, df)
t3 <- coef(summary(fit))[, "t value"][2]
```

1.6 (5 points)

Print a vector containing the values of `t1`, `t2`, and `t3`. What can you conclude from this? Why?

```
c(t1, t2, t3)
```

```
          t typewhite
9.684158  9.685650  9.685650
```

The similarity in the values of `t2` and `t3`, and the very close value of `t1`, suggest that regardless of the slight methodological differences, the statistical evidence pointing towards a certain hypothesis (such as the significance of a predictor in a regression model) is robust.

Question 2

25 points

Collinearity

2.1 (5 points)

Fit a linear regression model with all predictors against the response variable `quality`. Use the `broom::tidy()` function to print a summary of the fitted model. What can we conclude from the model summary?

```
library(broom)
# Fit the model
full_model <- lm(quality ~ ., data = df)
tidy_summary <- tidy(full_model)
# Print the summary using broom::tidy()
print(tidy_summary)
```

A tibble: 11 × 5

term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1 (Intercept)	57.5	9.33	6.17	7.44e-10
2 volatile_acidity	-1.61	0.0806	-20.0	4.07e-86
3 citric_acid	0.0272	0.0783	0.347	7.28e- 1
4 residual_sugar	0.0451	0.00416	10.8	3.64e-27
5 chlorides	-0.964	0.333	-2.90	3.78e- 3
6 total_sulfur_dioxide	-0.000329	0.000262	-1.25	2.10e- 1
7 density	-55.2	9.32	-5.92	3.34e- 9
8 pH	0.188	0.0661	2.85	4.38e- 3
9 sulphates	0.662	0.0758	8.73	3.21e-18
10 alcohol	0.277	0.0142	19.5	1.87e-82
11 typewhite	-0.386	0.0549	-7.02	2.39e-12

2.2 (10 points)

Fit two **simple** linear regression models using `lm()`: one with only `citric_acid` as the predictor, and another with only `total_sulfur_dioxide` as the predictor. In both models, use `quality` as the response variable. How does your model summary compare to the summary from the previous question?

```
# Model with citric_acid as the predictor
model_citric <- lm(quality ~ citric_acid, df)
summary(model_citric)
```

Call:

```
lm(formula = quality ~ citric_acid, data = df)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.9938	-0.7831	0.1552	0.2426	3.1963

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.65461	0.02602	217.343	<2e-16 ***
citric_acid	0.51398	0.07429	6.918	5e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8701 on 6495 degrees of freedom

Multiple R-squared: 0.007316, Adjusted R-squared: 0.007163

F-statistic: 47.87 on 1 and 6495 DF, p-value: 5.002e-12

```
# Model with total_sulfur_dioxide as the predictor
model_sulfur <- lm(quality ~ total_sulfur_dioxide, df)
summary(model_sulfur)
```

Call:

```
lm(formula = quality ~ total_sulfur_dioxide, data = df)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.8866	-0.7971	0.1658	0.2227	3.1965

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.8923848	0.0246717	238.831	< 2e-16 ***
total_sulfur_dioxide	-0.0006394	0.0001915	-3.338	0.000848 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8726 on 6495 degrees of freedom

Multiple R-squared: 0.001713, Adjusted R-squared: 0.001559

F-statistic: 11.14 on 1 and 6495 DF, p-value: 0.000848


```
model_sulfur <- ... # Insert your code here
```

2.3 (5 points)

Visualize the correlation matrix of all numeric columns in `df` using `corrplot()`

```
library(dplyr)      #
library(corrplot)

df %>%
  select(where(is.numeric)) %>% # Select only numeric columns
  cor() %>%                     # Compute the correlation matrix
  round(digits=2) %>%
  corrplot(diag=F)             # Generate the correlation plot
```

2.4 (5 points)

Compute the variance inflation factor (VIF) for each predictor in the full model using `vif()` function. What can we conclude from this?

```
library(car)
vif(full_model) %>% knitr::kable()
```

	x
volatile_acidity	2.103853
citric_acid	1.549248
residual_sugar	4.680035
chlorides	1.625065
total_sulfur_dioxide	2.628534
density	9.339357
pH	1.352005
sulphates	1.522809
alcohol	3.419849
type	6.694679

The VIF for residual_sugar, density and type are notably large. This shows the multi-collinearity in full_model

Question 3

40 points

Variable selection

3.1 (5 points)

Run a backward stepwise regression using a `full_model` object as the starting model. Store the final formula in an object called `backward_formula` using the built-in `formula()` function in R

```
library(MASS)
null_model <- lm(quality ~ 1, data = df)
full_model <- lm(quality ~ ., data = df)
backward_model <- stepAIC(full_model, direction = "backward")
```

Start: AIC=-3953.43

quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
total_sulfur_dioxide + density + pH + sulphates + alcohol +
type

	Df	Sum of Sq	RSS	AIC
- citric_acid	1	0.066	3523.6	-3955.3
- total_sulfur_dioxide	1	0.854	3524.4	-3953.9
<none>			3523.5	-3953.4
- pH	1	4.413	3527.9	-3947.3
- chlorides	1	4.559	3528.1	-3947.0
- density	1	19.054	3542.6	-3920.4
- type	1	26.794	3550.3	-3906.2
- sulphates	1	41.399	3564.9	-3879.5
- residual_sugar	1	63.881	3587.4	-3838.7
- alcohol	1	206.860	3730.4	-3584.8
- volatile_acidity	1	216.549	3740.0	-3567.9

Step: AIC=-3955.3

quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +
density + pH + sulphates + alcohol + type

	Df	Sum of Sq	RSS	AIC
- total_sulfur_dioxide	1	0.818	3524.4	-3955.8
<none>			3523.6	-3955.3
- chlorides	1	4.495	3528.1	-3949.0
- pH	1	4.536	3528.1	-3948.9
- density	1	20.794	3544.4	-3919.1

```

- type          1    26.943 3550.5 -3907.8
- sulphates     1    41.491 3565.1 -3881.2
- residual_sugar 1    67.371 3590.9 -3834.3
- alcohol       1   235.151 3758.7 -3537.6
- volatile_acidity 1  252.565 3776.1 -3507.5

```

Step: AIC=-3955.8

quality ~ volatile_acidity + residual_sugar + chlorides + density +
pH + sulphates + alcohol + type

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
- pH	1	4.295	3528.7	-3949.9
- chlorides	1	4.523	3528.9	-3949.5
- density	1	21.540	3545.9	-3918.2
- sulphates	1	40.711	3565.1	-3883.2
- type	1	43.664	3568.0	-3877.8
- residual_sugar	1	66.572	3591.0	-3836.2
- alcohol	1	244.545	3768.9	-3521.9
- volatile_acidity	1	256.695	3781.1	-3501.0

```
backward_formula <- formula(backward_model)
```

3.2 (5 points)

Run a forward stepwise regression using a `null_model` object as the starting model. Store the final formula in an object called `forward_formula` using the built-in `formula()` function in R

```
forward_model <- stepAIC(null_model, scope = list(lower = null_model, upper = full_model))
```

Start: AIC=-1760.04

quality ~ 1

	Df	Sum of Sq	RSS	AIC
+ alcohol	1	977.95	3975.7	-3186.9
+ density	1	463.41	4490.3	-2396.2
+ volatile_acidity	1	349.71	4604.0	-2233.7
+ chlorides	1	199.47	4754.2	-2025.1
+ type	1	70.53	4883.2	-1851.2
+ citric_acid	1	36.24	4917.4	-1805.7
+ total_sulfur_dioxide	1	8.48	4945.2	-1769.2
+ sulphates	1	7.34	4946.3	-1767.7
+ residual_sugar	1	6.77	4946.9	-1766.9
+ pH	1	1.88	4951.8	-1760.5
<none>			4953.7	-1760.0

Step: AIC=-3186.88

quality ~ alcohol

	Df	Sum of Sq	RSS	AIC
--	----	-----------	-----	-----

+ volatile_acidity	1	307.508	3668.2	-3707.9
+ residual_sugar	1	85.662	3890.1	-3326.4
+ type	1	54.335	3921.4	-3274.3
+ citric_acid	1	40.303	3935.4	-3251.1
+ chlorides	1	39.696	3936.0	-3250.1
+ total_sulfur_dioxide	1	31.346	3944.4	-3236.3
+ sulphates	1	7.859	3967.9	-3197.7
+ pH	1	5.938	3969.8	-3194.6
<none>			3975.7	-3186.9
+ density	1	0.005	3975.7	-3184.9

Step: AIC=-3707.89

quality ~ alcohol + volatile_acidity

	Df	Sum of Sq	RSS	AIC
+ sulphates	1	48.259	3620.0	-3791.9
+ density	1	38.704	3629.5	-3774.8
+ residual_sugar	1	29.751	3638.5	-3758.8
+ type	1	28.895	3639.3	-3757.3
+ total_sulfur_dioxide	1	5.619	3662.6	-3715.9
+ pH	1	5.533	3662.7	-3715.7
<none>			3668.2	-3707.9
+ chlorides	1	0.162	3668.1	-3706.2
+ citric_acid	1	0.099	3668.1	-3706.1

Step: AIC=-3791.94

quality ~ alcohol + volatile_acidity + sulphates

	Df	Sum of Sq	RSS	AIC
+ residual_sugar	1	43.989	3576.0	-3869.4
+ density	1	18.661	3601.3	-3823.5
+ type	1	6.012	3614.0	-3800.7
+ chlorides	1	4.988	3615.0	-3798.9
+ citric_acid	1	2.031	3617.9	-3793.6
+ pH	1	1.903	3618.1	-3793.4
<none>			3620.0	-3791.9
+ total_sulfur_dioxide	1	0.817	3619.2	-3791.4

Step: AIC=-3869.37

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar

	Df	Sum of Sq	RSS	AIC
+ type	1	20.7581	3555.2	-3905.2
+ total_sulfur_dioxide	1	13.3542	3562.6	-3891.7
+ pH	1	6.6430	3569.3	-3879.5
+ citric_acid	1	4.3384	3571.6	-3875.3
+ chlorides	1	1.8907	3574.1	-3870.8
<none>			3576.0	-3869.4
+ density	1	0.0071	3576.0	-3867.4

Step: AIC=-3905.19

```
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
  type
```

	Df	Sum of Sq	RSS	AIC
+ density	1	20.4623	3534.8	-3940.7
+ chlorides	1	6.6602	3548.6	-3915.4
+ citric_acid	1	5.2242	3550.0	-3912.7
+ pH	1	3.9477	3551.3	-3910.4
+ total_sulfur_dioxide	1	1.2539	3554.0	-3905.5
<none>			3555.2	-3905.2

Step: AIC=-3940.7

```
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
  type + density
```

	Df	Sum of Sq	RSS	AIC
+ chlorides	1	6.0826	3528.7	-3949.9
+ pH	1	5.8541	3528.9	-3949.5
<none>			3534.8	-3940.7
+ citric_acid	1	0.8471	3533.9	-3940.3
+ total_sulfur_dioxide	1	0.5646	3534.2	-3939.7

Step: AIC=-3949.89

```
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
  type + density + chlorides
```

	Df	Sum of Sq	RSS	AIC
+ pH	1	4.2945	3524.4	-3955.8
<none>			3528.7	-3949.9
+ total_sulfur_dioxide	1	0.5765	3528.1	-3948.9
+ citric_acid	1	0.2338	3528.4	-3948.3

Step: AIC=-3955.8

```
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
  type + density + chlorides + pH
```

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
+ total_sulfur_dioxide	1	0.81762	3523.6	-3955.3
+ citric_acid	1	0.02919	3524.4	-3953.9

```
forward_formula <- formula(forward_model)
```

3.3 (10 points)

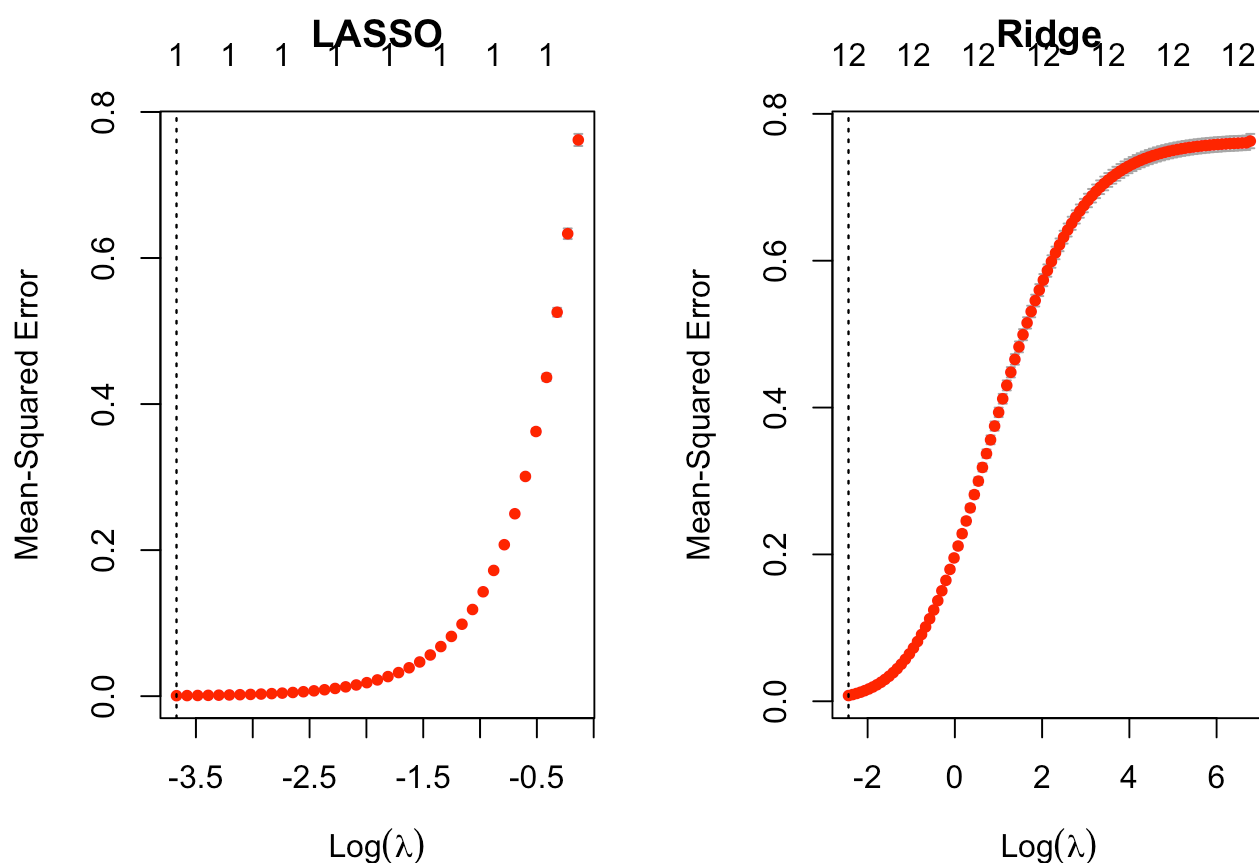
1. Create a `y` vector that contains the response variable (`quality`) from the `df` dataframe.
2. Create a design matrix `X` for the `full_model` object using the `make_model_matrix()` function provided in the Appendix.
3. Then, use the `cv.glmnet()` function to perform LASSO and Ridge regression with `X` and `y`.

```
library(glmnet)
y <- df$quality
X <- model.matrix(~ . -1, data = df) # Assuming 'df' is your dataset with 'quality' as on

lasso_model <- cv.glmnet(X, y, alpha = 1) # LASSO
ridge_model <- cv.glmnet(X, y, alpha = 0) # Ridge
```

Create side-by-side plots of the ridge and LASSO regression results. Interpret your main findings.

```
par(mfrow = c(1, 2)) # Setting up the plot area for two side-by-side plots
plot(lasso_model, main = "LASSO")
plot(ridge_model, main = "Ridge")
```



3.4 (5 points)

Print the coefficient values for LASSO regression at the `lambda.1se` value? What are the variables selected by LASSO?

Store the variable names with non-zero coefficients in `lasso_vars`, and create a formula object called `lasso_formula` using the `make_formula()` function provided in the Appendix.

```
lasso_coef <- coef(lasso_model, s = "lambda.1se")
lasso_vars <- rownames(lasso_coef)[lasso_coef[,1] != 0]
lasso_formula <- make_formula(lasso_vars[-1]) # Excluding the intercept
```

3.5 (5 points)

Print the coefficient values for ridge regression at the `lambda.1se` value? What are the variables selected here?

Store the variable names with non-zero coefficients in `ridge_vars`, and create a formula object called `ridge_formula` using the `make_formula()` function provided in the Appendix.

```
ridge_coef <- coef(ridge_model, s = "lambda.1se")
ridge_vars <- rownames(ridge_coef)[ridge_coef[,1] != 0]
ridge_formula <- make_formula(ridge_vars[-1]) # Excluding the intercept
```

3.6 (10 points)

What is the difference between stepwise selection, LASSO and ridge based on you analyses above?

LASSO tends to select a smaller number of variables compared to Ridge regression

Question 4

70 points

Variable selection

4.1 (5 points)

Excluding `quality` from `df` we have 10 possible predictors as the covariates. How many different models can we create using any subset of these 10 covariates as possible predictors? Justify your answer.

We can create $2^{10} - 1 = 1023$ different models. This is because each predictor can either be included or not included in a model,

From those 2^{10} combinations. We have to subtract 1 because we exclude the empty model the one with no predictors.

4.2 (20 points)

Store the names of the predictor variables (all columns except `quality`) in an object called `x_vars`.

```
# Install dplyr if it's not already installed
if (!require(dplyr)) install.packages("dplyr")
```

```
# Load dplyr, which will also load magrittr allowing you to use the pipe operator '%>%'
library(dplyr)

x_vars <- colnames(df %>% select(-quality))
if (!require(purrr)) install.packages("purrr")
library(purrr)
```

Use:

- the `combn()` function (built-in R function) and
- the `make_formula()` (provided in the Appendix)

to **generate all possible linear regression formulas** using the variables in `x_vars`. This is most optimally achieved using the `map()` function from the `purrr` package.

```
formulas <- map(
  1:length(x_vars),
  \(x) {
    vars <- combn(x_vars, x, simplify = FALSE)
    map(vars, make_formula)
  }
) %>% unlist()
```

If your code is right the following command should return something along the lines of:

```
sample(formulas, 4) %>% as.character()
# Output:
# [1] "quality ~ volatile_acidity + residual_sugar + density + pH + alcohol"
# [2] "quality ~ citric_acid"
# [3] "quality ~ volatile_acidity + citric_acid + residual_sugar + total_sulfur_dioxide +
# [4] "quality ~ citric_acid + chlorides + total_sulfur_dioxide + pH + alcohol + type"
```

4.3 (10 points)

Use `map()` and `lm()` to fit a linear regression model to each formula in `formulas`, using `df` as the data source. Use `broom::glance()` to extract the model summary statistics, and bind them together into a single tibble of summaries using the `bind_rows()` function from `dplyr`.

```
models <- map(formulas, \(f) lm(f, data = df))
summaries <- map(models, broom::glance) %>% bind_rows()
```

4.4 (5 points)

Extract the `adj.r.squared` values from `summaries` and use them to identify the formula with the **highest** adjusted R-squared value.

```
best_rsq <- summaries %>% filter(adj.r.squared == max(adj.r.squared))
```


Store resulting formula as a variable called `rsq_formula`.

```
rsq_formula <- formulas[which.max(summaries$adj.r.squared)]
```

4.5 (5 points)

Extract the `AIC` values from `summaries` and use them to identify the formula with the **lowest** AIC value.

```
best_aic <- summaries %>% filter(AIC == min(AIC))
```

Store resulting formula as a variable called `aic_formula`.

4.6 (15 points)

Combine all formulas shortlisted into a single vector called `final_formulas`.

```
aic_formula <- formulas[which.min(summaries$AIC)]
```

```
null_formula <- formula(null_model)
full_formula <- formula(full_model)

final_formulas <- c(
  null_formula,
  full_formula,
  backward_formula,
  forward_formula,
  lasso_formula,
  ridge_formula,
  rsq_formula,
  aic_formula
)
```

* Are ``aic_formula`` and ``rsq_formula`` the same? How do they differ from the formulas shortlisted in question 3?

* Which of these is more reliable? Why?

* If we had a dataset with \$10,000\$ columns, which of these methods would you consider for your analyses? Why?

4.7 (10 points)

Use ``map()`` and ``glance()`` to extract the ``sigma``, `adj.r.squared`, `AIC`, `df``, and ``p.value`` statistics for each model obtained from ``final_formulas``. Bind them together into a single data frame ``summary_table``. Summarize your main findings.

```
::: {.cell}
```

```
```{.r .cell-code}
print(colnames(df))
```

```
[1] "volatile_acidity" "citric_acid" "residual_sugar"
[4] "chlorides" "total_sulfur_dioxide" "density"
[7] "pH" "sulphates" "alcohol"
[10] "quality" "type"
```

```
print(final_formulas)
```

```
[[1]]
```

```
quality ~ 1
```

```
[[2]]
```

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
 total_sulfur_dioxide + density + pH + sulphates + alcohol +
 type
```

```
[[3]]
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + density +
 pH + sulphates + alcohol + type
```

```
[[4]]
```

```
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
 type + density + chlorides + pH
```

```
[[5]]
```

```
quality ~ quality
<environment: 0x12d1f55d8>
```

```
[[6]]
```

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
 total_sulfur_dioxide + density + pH + sulphates + alcohol +
 quality + typered + typewhite
<environment: 0x12d4e6db8>
```

```
[[7]]
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +
 density + pH + sulphates + alcohol + type
<environment: 0x11f7720a8>
```

```
[[8]]
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + density +
 pH + sulphates + alcohol + type
<environment: 0x11f7a44a8>
```

```
library(broom) # Ensure broom is loaded for using glance()

Modify the map function to include error handling
summary_table <- map(
 final_formulas,
 \(x) {
 tryCatch({
 model <- lm(x, data = df)
 glance(model) %>%
 select(sigma, adj.r.squared, AIC, df.residual, p.value) # Select the required st
 }, error = function(e) {
 message("Error in formula: ", deparse(x), "\nError Message: ", e$message)
 return(data.frame(sigma = NA, adj.r.squared = NA, AIC = NA, df.residual = NA, p.val
 })
 }
) %>%
bind_rows() # Combine all model summaries into one data frame
```

Warning in model.matrix.default(mt, mf, contrasts): the response appeared on the right-hand side and was dropped

Warning in model.matrix.default(mt, mf, contrasts): problem with term 1 in model.matrix: no columns are assigned

Error in formula: quality ~ volatile\_acidity + citric\_acid + residual\_sugar + chlorides + total\_sulfur\_dioxide + density + pH + sulphates + alcohol + quality + typered + typewhite

Error Message: object 'typered' not found

```
Adding names to the summary table for clarity
summary_table <- cbind(Formula = c("Null Formula", "Full Formula", "Backward Formula", "F
 "Ridge Formula", "Rsq Formula", "Aic Formula"), summar

Use knitr to create a better display table
summary_table %>% knitr::kable()
```

Formula	sigma	adj.r.squared	AIC	df.residual	p.value
Null Formula	0.8732553	0.0000000	16679.64	6496	NA
Full Formula	0.7370527	0.2876152	14486.26	6486	0
Backward Formula	0.7370314	0.2876563	14483.89	6488	0
Forward Formula	0.7370314	0.2876563	14483.89	6488	0
Lasso Formula	0.8732553	0.0000000	16679.64	6496	NA
Ridge Formula	NA	NA	NA	NA	NA

Formula	sigma	adj.r.squared	AIC	df.residual	p.value
Rsq Formula	0.7370027	0.2877118	14484.38	6487	0
Aic Formula	0.7370314	0.2876563	14483.89	6488	0

...

```

:::{.hidden unless-format="pdf"}
\pagebreak
:::

```

```

```

# Appendix

#### Convenience function for creating a formula object

The following function which takes as input a vector of column names `x` and outputs a `formula` object with `quality` as the response variable and the columns of `x` as the covariates.

```

::: {.cell}

```{r .cell-code}
make_formula <- function(x){
  as.formula(
    paste("quality ~ ", paste(x, collapse = " + "))
  )
}

:::

```

**For example the following code will
result in a formula object**

"quality ~ a + b + c"

```
make_formula(c("a", "b", "c"))
```

```
#### Convenience function for `glmnet`
```

The `make_model_matrix` function below takes a `formula` as input and outputs a **rescaled** model matrix `X` in a format amenable for `glmnet()`

```
::: {.cell}
```

```
```{.r .cell-code}
```

```
make_model_matrix <- function(formula){
 X <- model.matrix(formula, df)[, -1]
 cnames <- colnames(X)
 for(i in 1:ncol(X)){
 if(!cnames[i] == "typewhite"){
 X[, i] <- scale(X[, i])
 } else {
 colnames(X)[i] <- "type"
 }
 }
 return(X)
}
```

```
:::
```

## Session Information

Print your R session information using the following command

```
sessionInfo()
```

R version 4.3.3 (2024-02-29)

Platform: aarch64-apple-darwin20 (64-bit)

Running under: macOS Sonoma 14.4.1

Matrix products: default

BLAS: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRblas.0.dylib

LAPACK: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRlapack.dylib;

LAPACK version 3.11.0

locale:

[1] en\_US.UTF-8/en\_US.UTF-8/en\_US.UTF-8/C/en\_US.UTF-8/en\_US.UTF-8

time zone: America/New\_York

tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

other attached packages:

```
[1] broom_1.0.5 stringr_1.5.1 glmnet_4.1-8 Matrix_1.6-5
[5] car_3.1-2 carData_3.0-5 purrr_1.0.2 dplyr_1.1.4
[9] corrplot_0.92 tidyr_1.3.1 readr_2.1.5 MASS_7.3-60.0.1
```

loaded via a namespace (and not attached):

```
[1] utf8_1.2.4 generics_0.1.3 shape_1.4.6.1 stringi_1.8.3
[5] lattice_0.22-5 hms_1.1.3 digest_0.6.35 magrittr_2.0.3
[9] evaluate_0.23 grid_4.3.3 iterators_1.0.14 fastmap_1.1.1
[13] foreach_1.5.2 jsonlite_1.8.8 backports_1.4.1 survival_3.5-8
[17] fansi_1.0.6 codetools_0.2-19 abind_1.4-5 cli_3.6.2
[21] rlang_1.1.3 crayon_1.5.2 bit64_4.0.5 splines_4.3.3
[25] withr_3.0.0 tools_4.3.3 parallel_4.3.3 tzdb_0.4.0
[29] curl_5.2.1 vctrs_0.6.5 R6_2.5.1 lifecycle_1.0.4
[33] bit_4.0.5 vroom_1.6.5 pkgconfig_2.0.3 pillar_1.9.0
[37] glue_1.7.0 Rcpp_1.0.12 xfun_0.43 tibble_3.2.1
[41] tidyselect_1.2.1 rstudioapi_0.16.0 knitr_1.45 htmltools_0.5.8.1
[45] rmarkdown_2.26 compiler_4.3.3
```